



Homework #1

Carl Xi

(put your name above)

Note: This is an individual homework. Discussing this homework with your classmates is a violation of the Honor Code. If you borrow code from somewhere else, please add a comment in your code to make it clear what the source of the code is (e.g., a URL would sufficient). If you borrow code and you don't provide the source, it is a violation of the Honor Code.

Total grade: _____ out of ____150____ points

This homework has seven questions. Please answer them all and submit your assignment using Canvas. In particular, you need to submit:

- a) This file with your answers. Please transform the doc file to a PDF file. Then, submit the PDF File in the Assignment labeled as “Homework 1”***
- b) Your code and/or Rapidminer repositories for the hands-on exercise. Submit the (1) Python code in Jupyter format (2) the Rapidminer repositories and (3) the dataset(s) of the Assignment labeled as “Homework 1_Code & Repositories”***

1) (24 points) Choose the data technology (Q, H, U, or S) that is most appropriate for each of the following business questions/scenarios.

Q – SQL Querying
H – Statistical Hypothesis Testing
U – Unsupervised Data Mining/Pattern Finding
S – Supervised Data Mining

- a) **_H_** For my on-line advertising the decision tree model yields a response rate of 0.5% and the old, manual targeting model yields 0.3%. Is the decision tree model really better?
- b) **_Q_** I want to know which of my customers are the most profitable.
- c) **_Q_** I need to get data on all my on-line customers who were exposed to the special offer, including their registration data, their past purchases, and whether or not they purchased in the 15 days following the exposure.
- d) **_U_** I would like to segment my customers into groups based on their demographics and prior purchase activity. I am not focusing on improving a particular task, but would like to generate ideas.
- e) **_S_** I have a budget to target 10,000 existing customers with a special offer. I would like to identify those customers most likely to respond to the special offer.
- f) **_U_** I want to know what characteristics differentiate my most profitable customers.

2) (16 points) Label each case as describing either data mining (DM), or the use of the results of data mining (Use).

- g) **_Use_** Choose customers who are most likely to respond to an on-line ad.
- h) **_DM_** Discover rules that indicate when an account has been defrauded.
- i) **_DM_** Find patterns indicating what customer behavior is more likely to lead to response to an on-line ad.
- j) **_Use_** Estimate probability of default for a credit application.

3) (15 points) MTC (MegaTelCo) has decided to use supervised learning to address its problem of churn in its wireless phone business. As a consultant to MTC, you realize that a main task in the business understanding/data understanding phases of the data mining process is to define the target variable. In one or two sentences, please suggest a definition for the target variable. Be as precise as possible—someone else will be implementing your suggestion. (Remember: it should make sense from a business point of view, and it should be reasonable that MTC would have data available to know the value of the target variable for historical customers.)

The target variable is essentially the attribute within a dataset that you are trying to predict with the rest of the attributes in your available data. In this case, the target variable would be the churn variable of MTC’s wireless phone business dataset. We will be trying to create a predictive model (a formula) that estimates whether a customer will churn within 3 months of contract expiration (the target variable) given a specific set of attributes. In this case, using class probability estimation methods will probably yield better results than binary probability estimation methods.

4) (20 points) A predictive model has been applied to a test dataset and has classified 87 records as fraudulent (31 correctly so) and 953 as non-fraudulent (919 correctly so).

- **Present the confusion matrix for this scenario.**

Predicted\Actual	Positive (Fraudulent)	Negative (Non-Fraudulent)	Total
Positive (Fraudulent)	31 (True Positives [TP])	56 (False Positives [FP])	87
Negative (Non-Fraudulent)	34 (False Negatives [FN])	919 (True Negatives [TN])	953

Total	65	975	1040
--------------	----	-----	------

- Calculate the error rate and accuracy rate.

Accuracy = $(TP+TN)/(TP+FP+FN+TN) = (31+919)/(31+56+34+919) = 0.9135$ or 91.35%

Error rate = $1 - \text{Accuracy} = 1 - 0.9135 = 0.0865$ or 8.65%

- Calculate the precision, recall, and f-measure values for each of the two outcome classes (i.e., fraudulent and non-fraudulent records);

Precision _{Positive} = $TP / (TP + FP)$	Precision _{Negative} = $TN / (TN + FN)$
Recall _{Positive} = $TP / (TP + FN)$	Recall _{Negative} = $TN / (TN + FP)$
F-Measure _{Positive} = $2 * \frac{\text{Precision[Positive]} * \text{Recall[Positive]}}{\text{Precision[Positive]} + \text{Recall[Positive]}}$	F-Measure _{Negative} = $2 * \frac{\text{Precision[Negative]} * \text{Recall[Negative]}}{\text{Precision[Negative]} + \text{Recall[Negative]}}$

Positive = Fraudulent, Negative = Non-Fraudulent

Precision _{Positive} = $31/(31+56) = \frac{31}{87} \approx 0.3563$	Precision _{Negative} = $919/(919+34) = \frac{919}{953} \approx 0.9643$
Recall _{Positive} = $31/(31+34) = \frac{31}{65} \approx 0.4769$	Recall _{Negative} = $919/(919+56) = \frac{919}{975} \approx 0.9426$
F-Measure _{Positive} = $2 * \frac{0.3563 * 0.4769}{0.3563 + 0.4769} \approx 0.4079$	F-Measure _{Negative} = $2 * \frac{0.9643 * 0.9426}{0.9643 + 0.9426} \approx 0.9533$

- Also, calculate the accuracy rate that would be achieved by naïve (majority) rule on this data.

The naïve rule basically states that all data records will be marked as ‘non-fraudulent’, or that they belong to the majority (most prevalent class).

Naïve accuracy rate = $56+919 = \frac{(56+919)}{\text{total}} = \frac{(56+919)}{(31+34+56+919)} = 0.9375$

5) (50 points) [Implement this exercise with both RapidMiner (20 points) and Python (30 points)] Use the decision tree classification technique on the *HWI* dataset. This dataset is provided on the course website and contains data about consumers and their decisions to terminate a contract (i.e., consumer churn problem).

Data description:

Col. Var. Name Var. Description

1	revenue	Mean monthly revenue in dollars
2	outcalls	Mean number of outbound voice calls
3	incalls	Mean number of inbound voice calls
4	months	Months in Service
5	eqpdays	Number of days the customer has had his/her current equipment
6	webcap	Handset is web capable
7	marryes	Married (1=Yes; 0=No)
8	travel	Has traveled to non-US country (1=Yes; 0=No)
9	pcown	Owens a personal computer (1=Yes; 0=No)
10	credited	Possesses a credit card (1=Yes; 0=No)
11	retcalls	Number of calls previously made to retention team
12	churndep	Did the customer churn (1=Yes; 0=No)

Build a decision tree model that predicts whether a consumer will terminate his/her contract. In particular, I would like for you to create a decision tree using entropy with no max depth. Explore how well the decision trees perform for several different parameter values (e.g., for different splitting criteria). Interpret the model (decision tree) that provides the best predictive performance. Some possible issues / hints to think about: using training vs. test datasets. Present a brief overview of your predictive modeling process, explorations, and discuss your results. Make sure you present information about the model “goodness” (please report the confusion matrix, predictive accuracy, classification error, precision, recall, f-measure).

Present a brief overview of your predictive modeling process. That is, you need to lay out the steps you have taken in order to build and evaluate the decision tree model. For instance, how did you explore the data set before you built the model? Write this report in a way that the upper level management of the team would understand what you are doing. Why is the decision tree an appropriate model for this problem? How can we evaluate the predictive ability of the decision tree? If you build decision trees with different splitting criteria, which decision tree would you prefer to use in practice?

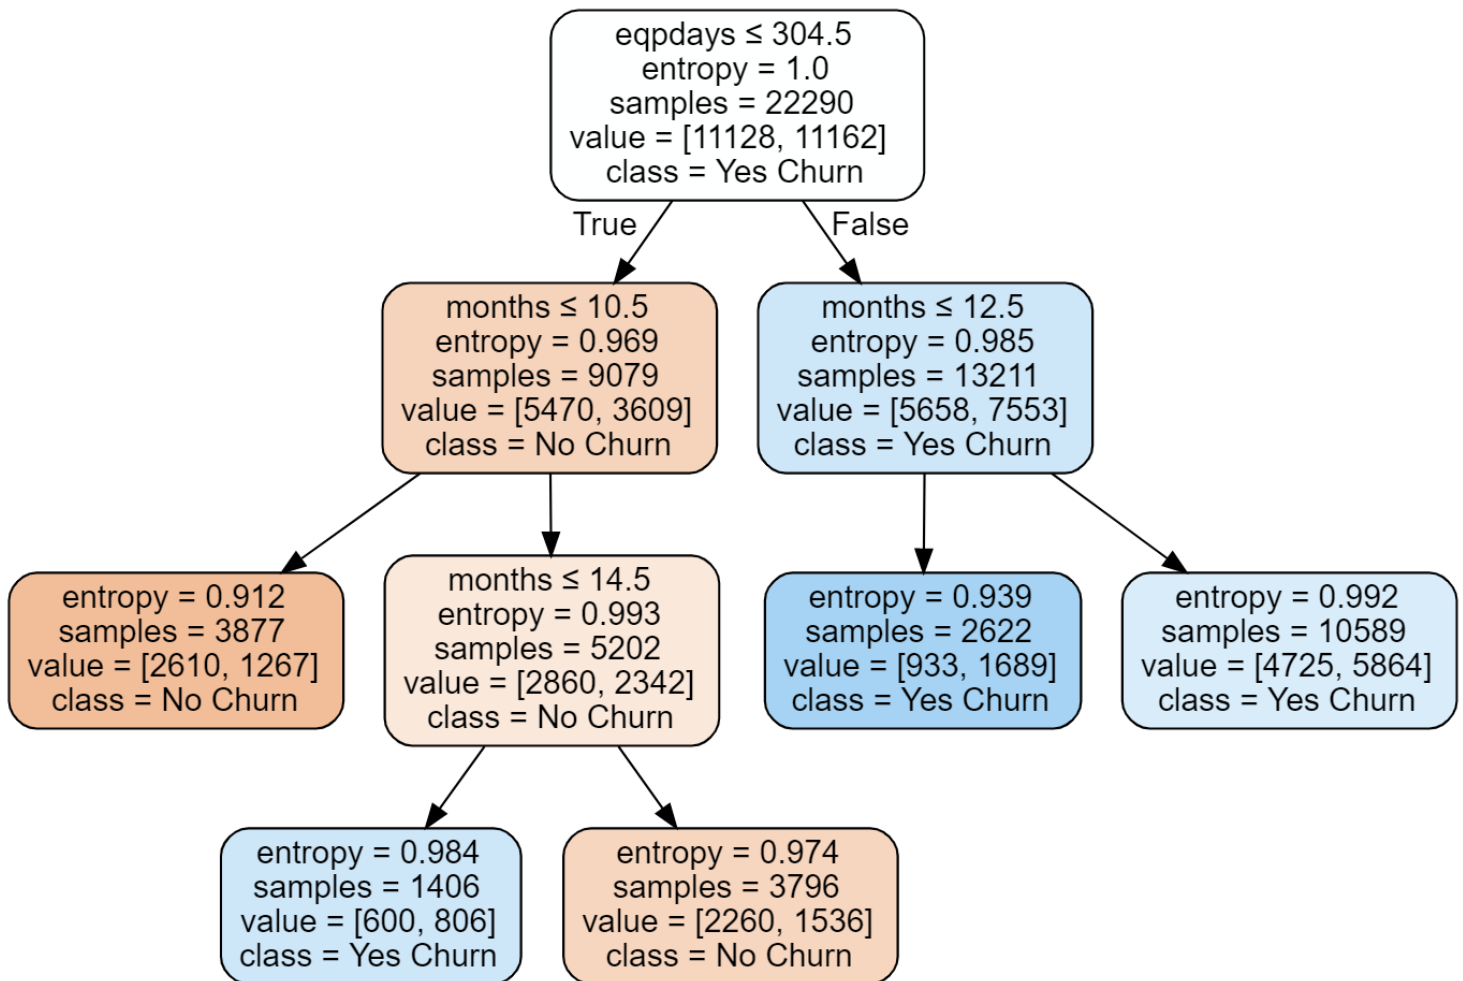
Business problem & data set overview:

Following our answer in question 3, we are essentially trying to build a model that predicts whether a customer will churn given a set of variables. We can immediately identify this as a supervised classification problem as we have a specific purpose and existing data on the target variable. Benefits of decision trees for this problem are discussed at the bottom of this problem.

In terms of our data, we can see that 4 of our variables are 0/1 or binary, while the rest are all floating numbers. Our target variable is also 0/1 or binary.

[Python Model]

First, here is my decision tree in python:



Here are the variables I used in my DecisionTreeClassifier call:

```
DecisionTreeClassifier(criterion='entropy', random_state= 42, max_depth = None, min_samples_split = 0.081,  
min_impurity_decrease = 0.002)
```

- Note that we have set max_depth to 'None', but that other variables like min_samples_split and min_impurity_decrease will still affect the total number of branches.

Here are the results of my best decision tree:

```
In [11]: ##### Confusion Matrix #####

# Compute confusion matrix to evaluate the accuracy of a classification
cnf_matrix = confusion_matrix(y_test, y_pred)
np.set_printoptions(precision=2)

print("Confusion Matrix:")
print(confusion_matrix(y_test,y_pred))
print()
print("Summary Statistics")
print(classification_report(y_test,y_pred))
print()
print("Accuracy:", metrics.accuracy_score(y_test,y_pred))
```

Confusion Matrix:

```
[[2123 2755]
 [1132 3544]]
```

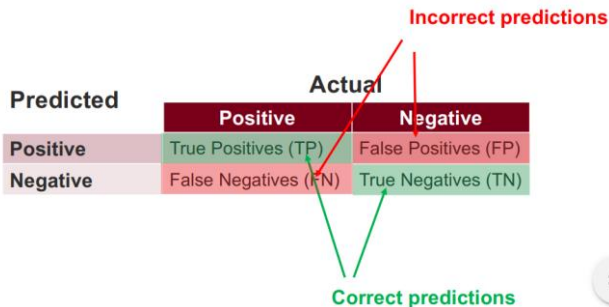
Summary Statistics

	precision	recall	f1-score	support
0	0.65	0.44	0.52	4878
1	0.56	0.76	0.65	4676
micro avg	0.59	0.59	0.59	9554
macro avg	0.61	0.60	0.58	9554
weighted avg	0.61	0.59	0.58	9554

Accuracy: 0.5931546996022609

A quick review on the definitions of the statistics:

- Predictive accuracy: % of correct classifications out of total.
- Classification error: % of wrong classifications out of total.
- An accuracy of 0.5932 is very close to the accuracy I achieved with my RapidMiner model (which I will discuss below). The confusion matrix can be interpreted using the diagram below:



We can see that there are a sizable number of False Positives and False Negatives, but in-line with our 0.5932 accuracy. Our f-measure of 0.52 and 0.65 for no churn and churn respectively show us that our decision tree performs above average on a scale of 0 to 1.

Before I got my final decision tree, I played around with the different variables in the DecisionTreeClassifier call. I have produced about 10 different decision trees, but to save space I have only included 3 in my writeup. With my final tree excluded, the other two notable ones are below:

1. With nothing but random state, the call prints out a scary big and long tree that is effectively illegible.
2. My next significant improvement to my model is when I set the minimal sample split to 0.08. It produced a legible model that gets the job done, but I was still unhappy with how big it is.


```
In [1]: # To write a Python 2/3 compatible codebase, the first step is to add this line to the top of each module
from __future__ import division, print_function, unicode_literals

##### Imports #####
# Numpy is the fundamental package for scientific computing with Python.
# SciPy (pronounced "Sigh Pie") is an open source Python library used for scientific computing and technical computing.
# Os module provides a portable way of using operating system dependent functionality.
import numpy as np # np is an alias pointing to numpy
import scipy as sp # sp is an alias pointing to scipy
import pandas as pd # pd is an alias point to pandas

from sklearn import metrics, tree
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split
from sklearn.tree import export_graphviz
# If you don't have graphviz package, you need to install it https://anaconda.org/anaconda/graphviz
# How to install Graphviz with Anaconda 1
# conda install -c anaconda graphviz
from IPython.display import Image
import matplotlib.pyplot as plt
!pip install graphviz
import itertools
import graphviz
import os

# Seed the generator to make this notebook's output stable across runs
np.random.seed(42)

# To plot pretty figures
%matplotlib inline
import matplotlib
import matplotlib.pyplot as plt #pyplot is matplotlib's plotting framework https://matplotlib.org/users/pyplot_tutorial.html
```

I then imported the dataset in its raw state

```
##### Imports #####
hw1data = pd.read_csv("HW1_Data.csv")
```

I know from my work in RapidMiner that there are some outliers and bad data, so I wanted to examine the data and eliminate them:

```
In [2]: #Looking at how big our dataset is
hw1data.shape
```

Out[2]: (31891, 12)

```
In [3]: #taking a peek at how the data looks like
hw1data.head()
```

Out[3]:

	revenue	outcalls	incalls	months	eqpdays	webcap	marryyes	travel	pcown	creditcd	retcalls	churndep
0	83.53	20.00	1.0	31	745	1	0	0	0	0	4	1
1	29.99	0.00	0.0	52	1441	0	0	0	1	1	3	1
2	37.75	2.67	0.0	25	572	0	0	0	1	1	3	1
3	5.25	0.00	0.0	45	1354	0	0	0	0	0	2	1
4	42.71	8.67	0.0	27	224	1	0	0	0	0	3	1

```
In [4]: #Looking at any potential problems/outliers
hw1data.describe()
```

Out[4]:

	revenue	outcalls	incalls	months	eqpdays	webcap	marryyes	travel	pcown	creditcd	re
count	31891.000000	31891.000000	31891.000000	31891.000000	31891.000000	31891.000000	31891.000000	31891.000000	31891.000000	31891.000000	31891.000000
mean	58.665179	24.951385	8.065277	18.761908	391.222633	0.894704	0.363175	0.057163	0.184817	0.676931	0.000000
std	44.163859	34.790147	16.610589	9.548019	254.998478	0.306939	0.480922	0.232158	0.388155	0.467656	0.200000
min	-5.860000	0.000000	0.000000	6.000000	-5.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	33.450000	3.000000	0.000000	11.000000	212.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	48.380000	13.330000	2.000000	17.000000	341.000000	1.000000	0.000000	0.000000	0.000000	1.000000	0.000000
75%	71.040000	33.330000	9.000000	24.000000	530.000000	1.000000	1.000000	0.000000	0.000000	1.000000	0.000000
max	861.110000	610.330000	404.000000	60.000000	1812.000000	1.000000	1.000000	1.000000	1.000000	1.000000	4.000000

I noticed some problems, such as how the min for revenue and eqpdays were negative and removed these bad data points from my dataset.


```
In [5]: #data cleaning
baddata = hw1data[(hw1data['revenue']<0) | (hw1data['eqpdays']<0)].index
hw1data.drop(baddata, inplace=True)
```

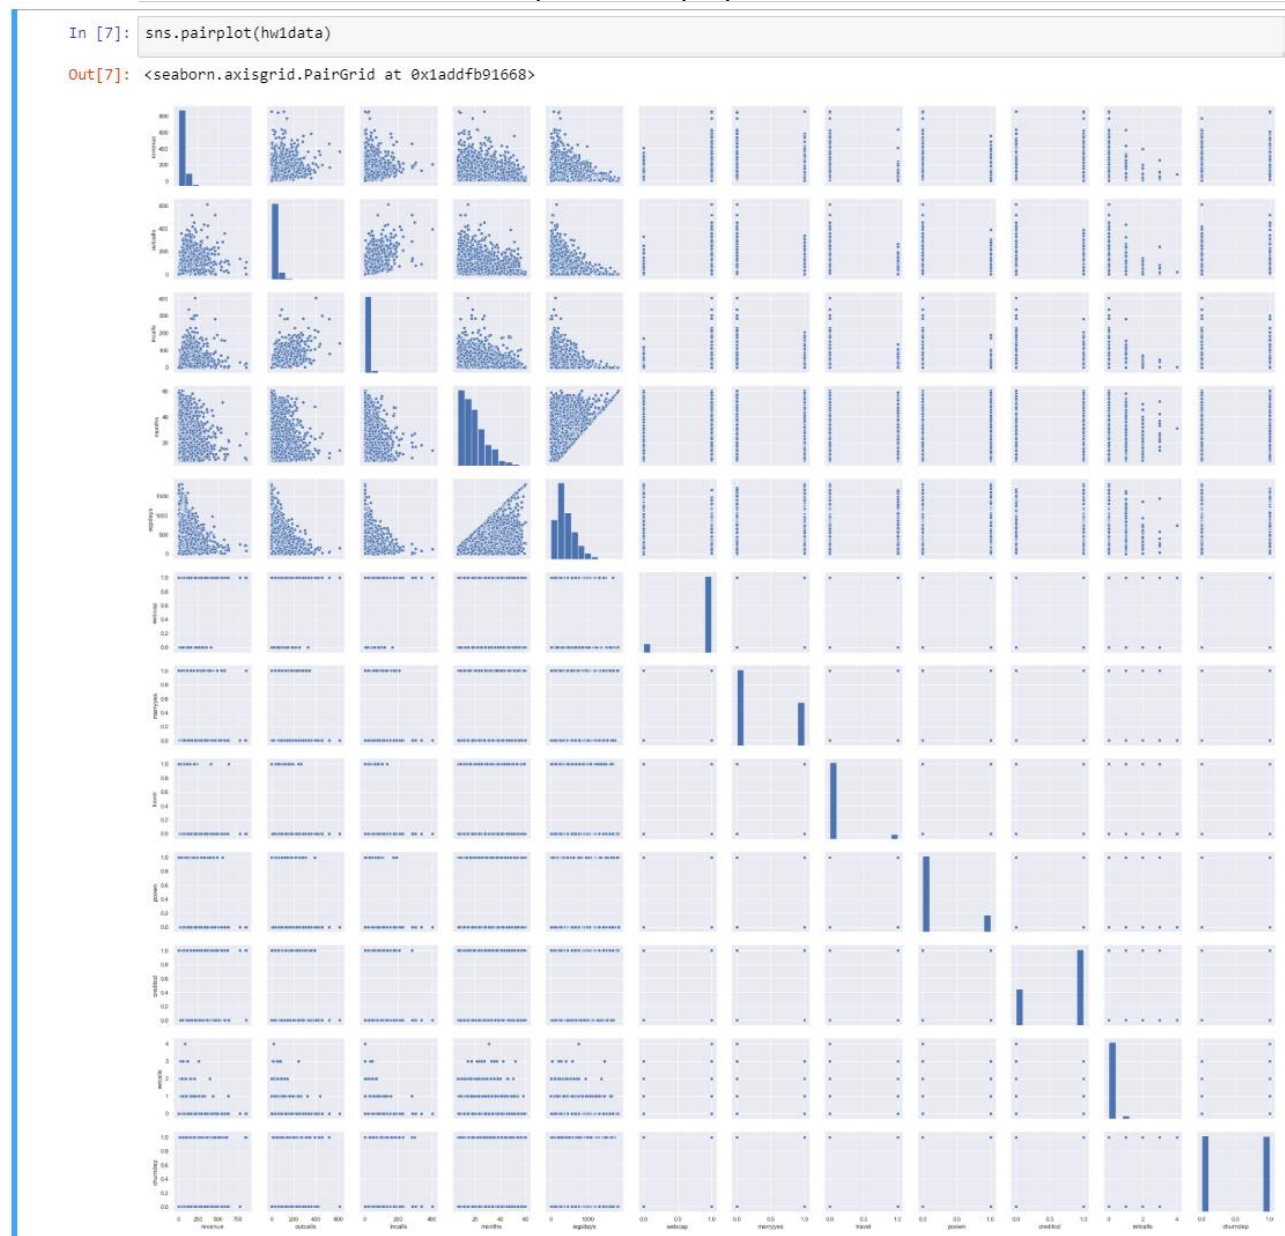
Then I imported the packages for decision tree making, and isolated my target variable from the dataset:

```
In [6]: from sklearn.tree import DecisionTreeClassifier # The sklearn.tree module includes decision tree-based models for
# classification and regression
# Documentation for decision Tree Classifier
# http://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html

# Seaborn is a Python data visualization library based on matplotlib.
# Seaborn documentation can be found here https://seaborn.pydata.org/generated/seaborn.set.html
import seaborn as sns # sns is an alias pointing to seaborn
sns.set(color_codes=True) #Set aesthetic parameters in one step. Remaps the shorthand color codes (e.g. "b", "g", "r", etc.) to
from scipy import stats #Documentation stats package of scipy https://docs.scipy.org/doc/scipy/reference/stats.html#module-scipy

##### Isolating the Target Variable #####
# Retrieving Attributes
X = hw1data.iloc[:, :-1]
# Retriving Target Variable
y = hw1data.iloc[:, -1]
```

With the data split and cleaned, I decided to take a final look at the trends in my data to see if there are patterns that might influence our model, as well as cross-variable relationships. Here is a pairplot matrix :



From the plot, I noticed that revenue, outcalls and incalls are heavily skewed towards to the left (smaller values). Additionally, most entries have web capable handsets, have traveled to non-US countries, owns a personal computer and previously did not call the retention team. Our target variable (churndep) does not have obvious correlations with any other variables in our dataset.

Explore how well the decision trees perform for several different parameter values (e.g., for different splitting criteria). Interpret the model (decision tree) that provides the best predictive performance.

Why is the decision tree an appropriate model for this problem? How can we evaluate the predictive ability of the decision tree?

I then split my data into 70% train and 30% test, and ran my model:

```
In [8]: ##### Split the Data #####

# Split the data into a training set and a test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
# Note: Always a good idea to shuffle the dataset before you split it into training and testing
# train_test_split performs shuffling by default

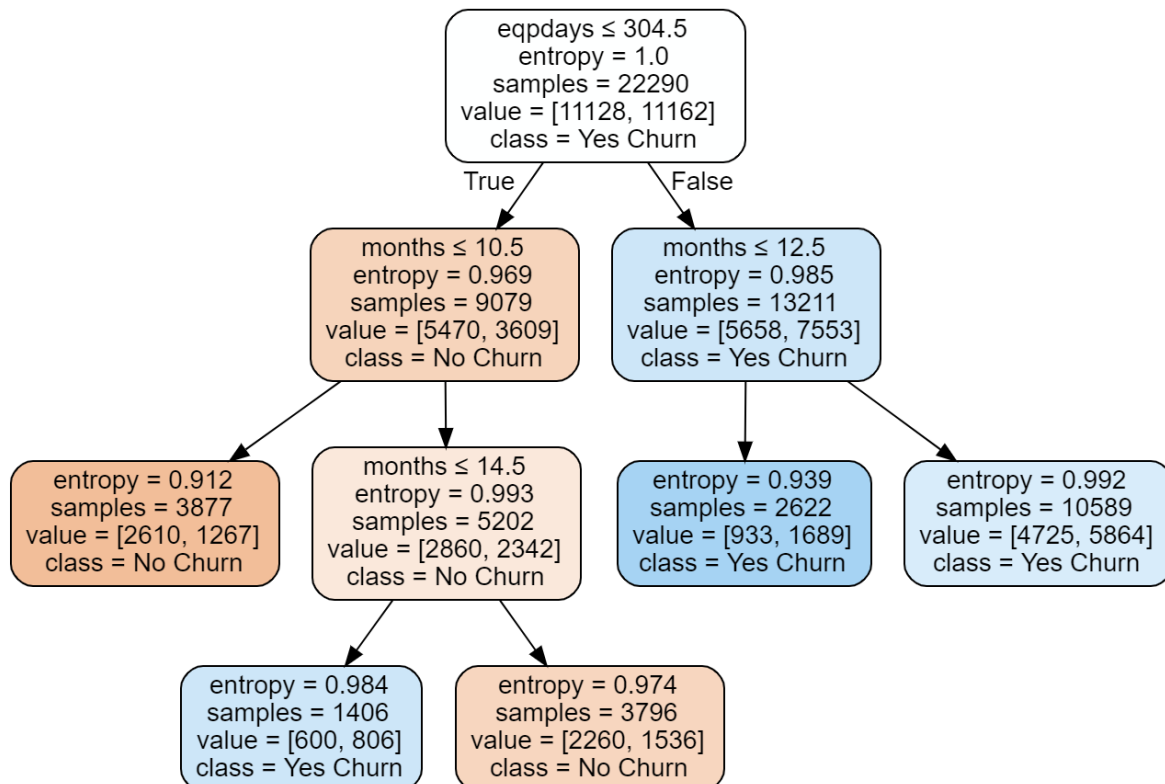
In [9]: ##### Model Training #####
##### Build Model & Apply it to the Test Set #####
#Build the decision tree
clf3 = DecisionTreeClassifier(criterion='entropy', random_state= 42,max_depth = None, min_samples_split=0.081, min_impurity_dec
# "clf3.fit(X_train, y_train)" fits the model and then
# ".predict(X_test)" makes predictions based on the test set
y_pred = clf3.fit(X_train, y_train).predict(X_test)
```

With my model successfully ran, I printed my tree to examine it visually:

```
In [10]: ##### Visualization of Decision Tree #####
banana = tree.export_graphviz( clf3,
    out_file=None, #image_path("HW1.dot")
    feature_names=list(X.columns),
    class_names=["No Churn","Yes Churn"],
    rounded=True,
    filled=True,
    special_characters=True)

graph = graphviz.Source(banana)
graph
```

Out[10]:



Lastly, I printed out all my evaluation metrics to see how well my model performs. This includes the confusion matrix, f-score, precision, and accuracy. I have discussed the meaning of these results in more detail at the beginning of this question.

```
In [11]: ##### Confusion Matrix #####

# Compute confusion matrix to evaluate the accuracy of a classification
cnf_matrix = confusion_matrix(y_test, y_pred)
np.set_printoptions(precision=2)

print("Confusion Matrix:")
print(confusion_matrix(y_test,y_pred))
print()
print("Summary Statistics")
print(classification_report(y_test,y_pred))
print()
print("Accuracy:", metrics.accuracy_score(y_test,y_pred))
```

Confusion Matrix:

```
[[2123 2755]
 [1132 3544]]
```

Summary Statistics

	precision	recall	f1-score	support
0	0.65	0.44	0.52	4878
1	0.56	0.76	0.65	4676
micro avg	0.59	0.59	0.59	9554
macro avg	0.61	0.60	0.58	9554
weighted avg	0.61	0.59	0.58	9554

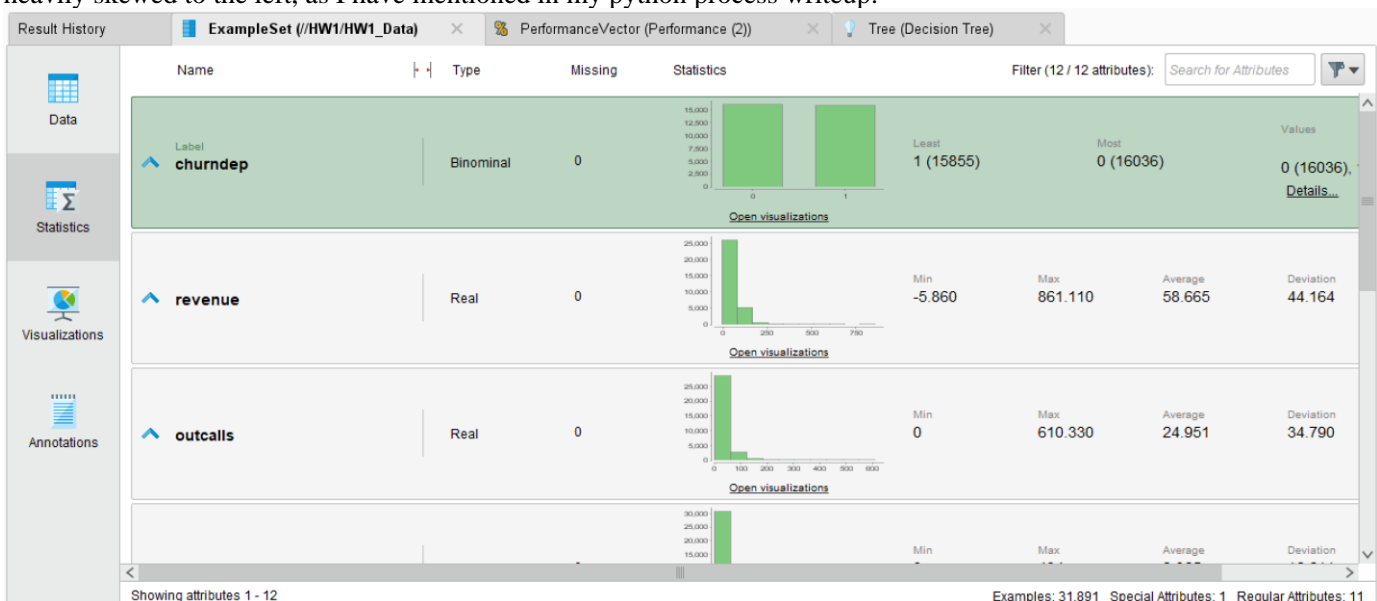
Accuracy: 0.5931546996022609

[RapidMiner Process]

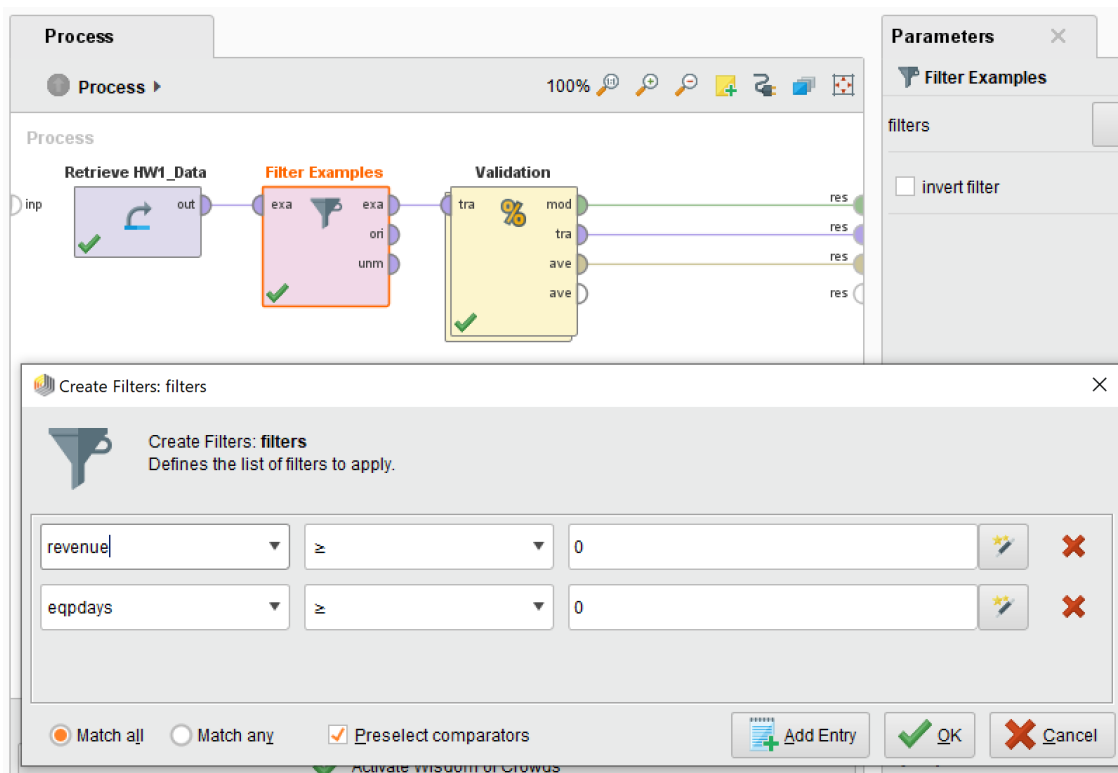
I first created a new repository in RapidMiner and moved my data into it. I then saved a copy of the in class-exercise to build off of for my RapidMiner model:



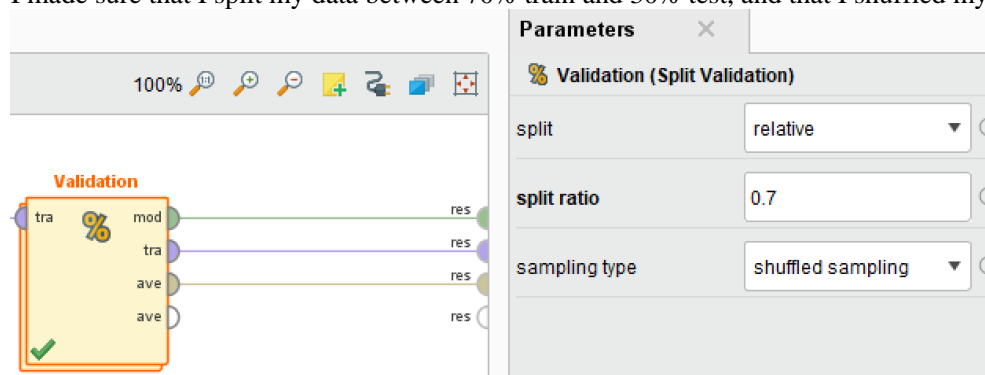
I then looked at my variables and tried to identify any concerning datapoints. For example, revenue and eqpdays have negative numbers for their mins. I kept this in mind as I took a look at the histograms of all my variables. I noticed that many of them are heavily skewed to the left, as I have mentioned in my python process writeup:



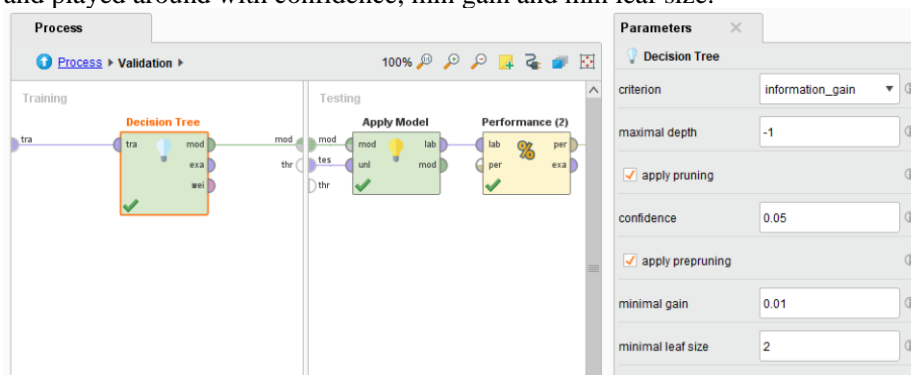
Going back to my process map, I added a filter to remove all the bad data points mentioned in my previous step:



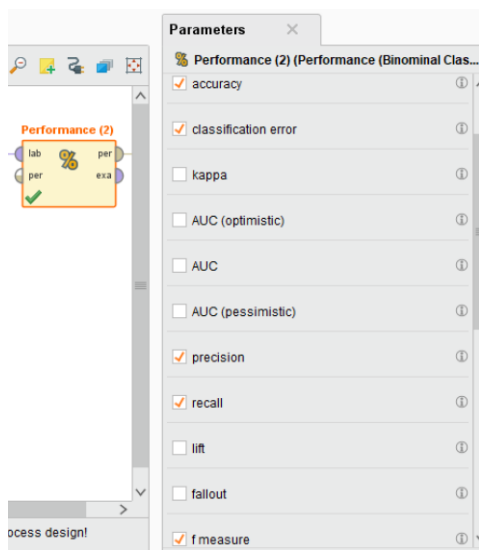
I made sure that I split my data between 70% train and 30% test, and that I shuffled my samples.



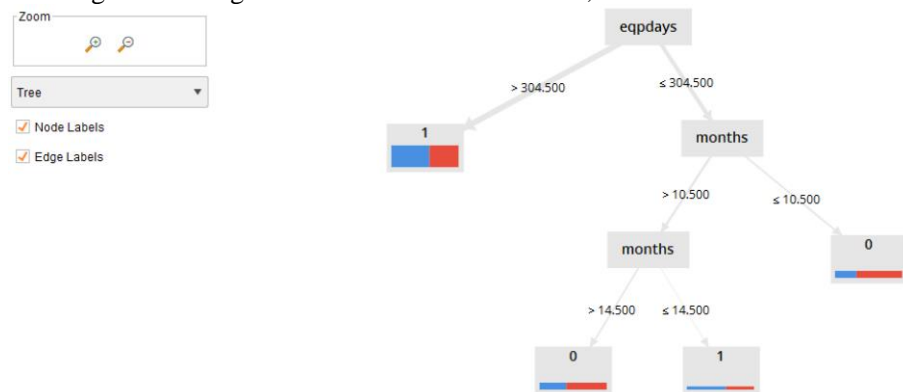
In the validation module, I made sure to change my criterion to information gain, max depth to -1 (this basically means no limit), and played around with confidence, min gain and min leaf size.



I then made sure that I got the output statistics on my model that I wanted by selecting them under the performance module:



Running the model gave us our model decision tree, and its statistics below:



Tree

```
eqpdays > 304.500: 1 {1=10729, 0=8165}
eqpdays ≤ 304.500
| months > 10.500
| | months > 14.500: 0 {1=2173, 0=3254}
| | months ≤ 14.500: 1 {1=1174, 0=848}
| months ≤ 10.500: 0 {1=1762, 0=3739}
```

accuracy	accuracy: 59.73%		
classification error			
precision			
recall			
f measure			
false positive			
false negative			
true positive			
true negative			
sensitivity			

	true 1	true 0	class precision
pred. 1	3673	2720	57.45%
pred. 0	1127	2033	64.34%
class recall	78.52%	42.77%	

All the individual statistics are in their own tab in the screenshot above. To save space, I have just typed them out below:

Classification error: 40.27%

Precision: 64.34% (positive)

Recall: 42.77% (positive)

f-measure: 51.38 (positive)

All True negative, True positive, false positive, sales negative values can be seen in the confusion matrix in the screenshot above.

Sensitivity: 42.77% (positive)

Similar to the summary statistics from my decision tree from the python model, our model is 'above' average given our f-measure, recall and precision values. The decision tree is an appropriate model for this problem because it helps us maximize information gain and is relatively efficient at classification problems. We can evaluate the predictive ability of the decision tree mostly by looking at the accuracy. Both my RapidMiner and Python models have an accuracy of around 59%, meaning that 59% of the time the model can correctly identify whether someone will churn or not. Given churn/not churn is 50/50, it's a small but still noticeable improvement of 9%. In practice, I would prefer making a decision tree with parameters that are the most effective at maximizing information gain and accuracy, while keeping the tree concise. When we did not have an additional parameter in the python model, the resulting decision tree looked like something straight out of a horror film. Also, upper management will want an answer/model that's the easiest to understand while being the most effective. An overly complex one will not satisfy the first

requirement.

6) (20 points) Is a node's entropy generally lower or greater than its parent's? Is it ever possible for a node's entropy to be higher than its parent's entropy? Please justify your answer. Be precise but concise.

While a node's entropy is generally lower than the entropy of its parent, it is possible for a node's entropy to be higher than its parent's. This is dependent on how you split the parent node.

However, with the goal of increasing the purity of the data set in mind, typically a node's entropy should be lower than that of the parent's as smaller entropy leads to higher dataset purity.

Here's an example:

	Parent Node (10x, 5y)	
Child Node 1 (5x, 4y)		Child Node 2 (5x, 1y)

Parent Entropy = $-(5/15)*\log_2(5/15)-(10/15)*\log_2(10/15) = 0.9183$

Child 1 Entropy = $-(5/9)*\log_2(5/9)-(4/9)*\log_2(4/9) = 0.9911$

Child 2 Entropy = $-(5/6)*\log_2(5/6)-(1/6)*\log_2(1/6) = 0.6500$

Total information gain: $0.9183-(9/15)*0.9911-(6/15)*0.6500 = 0.06364$

Despite there being information gain, one of the two child nodes have a higher entropy than the parent node.

7) (5 points) What are the differences between supervised and unsupervised methods in machine learning? Is the decision tree algorithm a supervised or an unsupervised method? Be precise but concise.

Usually, supervised data mining results produces a model that predicts a quantifiable target variable of interest. Unsupervised methods are conducted with a specific target or purpose in mind while supervised methods do not. Supervised methods also require data on the target variable. Also, sometimes the results of unsupervised data mining methods are not useful or meaningful at all for any prediction purpose while supervised results are typically useful to a certain degree.

Decision tree algorithm falls under supervised methods because it is used for classification problems.